

# Concept learning

- ▶ Inferring a boolean-valued function from training examples of its input and output.
- ▶ Concept learning is also known as binary classification.
- ▶ Example: "EnjoySport"

Suppose want to learn target concept days on which Fred enjoys his favorite *water sport*

## Concept learning ...

- Input is a set of examples, one per day
  - describing the day in terms of a set of *attributes*
  - indicating (yes/no) whether Fred enjoyed his sport that day

Example	Sky	Temp	Humid	Wind	Water	Forecast	EnjoySport
1	Sunny	Warm	Normal	Strong	Warm	Same	Yes
2	Sunny	Warm	High	Strong	Warm	Same	Yes
3	Rainy	Cold	High	Strong	Warm	Change	No
4	Sunny	Warm	High	Strong	Cool	Change	Yes

- Task: learn to predict the value of *EnjoySport* for an arbitrary day, given values of other attributes

# Concept learning ...

- Suppose hypotheses take form of conjunctions of constraints on instance attributes – e.g. specify allowed values of:  
*Sky, Temp, Humid, Wind, Water, Forecast*  
and suppose constraints take one of three forms:
  - ? – any value acceptable
  - Specified\_value – specific value required, e.g. *Warm* for the *Temp* attribute
  - $\emptyset$  – no value acceptable
- Then, hypotheses can be represented as vectors of such constraints.  
E.g.  $\langle ?, Cold, High, ?, ?, ? \rangle$  – represents hypothesis  
*Fred enjoys sport only on cold days with high humidity*

# Concept learning ...

- Can describe the concept learning setting as follows.  
Given:
  - $X$  a set of *instances* over which the concept is to be defined (each represented, e.g., as a vector of attribute values)
  - a *target function* or *concept* to be learned:
 
$$c : X \rightarrow \{0, 1\}$$
  - a set  $D$  of *training examples*, each of the form  $\langle x, c(x) \rangle$  where  $x \in X$  and  $c(x)$  is the target concept value for  $x$   
(Note: instances in  $D$  for which  $c(x) = 1$  are called *positive examples*, those for which  $c(x) = 0$  are *negative examples*)

# Concept learning ...

Find:

- a *hypothesis*, or estimate, of  $c$ .

I.e. supposing  $H$  is the set of all hypotheses, find  $h \in H$ , where  $h : X \rightarrow \{0, 1\}$  such that

$$h(x) = c(x) \text{ for all } x \in X$$

- Thus, concept learning can be viewed as **search** over the space of hypotheses, as defined by the hypothesis representation.
- Note: we want to find  $h$  identical to  $c$  over all of  $X$  **but**, only have information about  $c$  for training examples  $D$ .
- Inductive learning algorithms can only guarantee that hypotheses fit *training data*.

# Concept learning ...

**Inductive Learning Hypothesis** Any hypothesis found to approximate the target function well over a sufficiently large set of training examples will also approximate the target function well over other unobserved examples.

## General-to-Specific Ordering of Hypotheses

- **Problem:** How should we search this space to find the target concept?  
**A Solution:** Start with the most specific hypothesis and, considering each training example in turn, generalise towards the most general hypothesis, stopping at the first hypothesis that ‘covers’ the training examples (FIND-S)

# Concept learning ...

- the most specific hypothesis:

$\langle 0, 0, 0, 0, 0, 0 \rangle$  (*Fred enjoys sport on no day*)

- the most general hypothesis:

$\langle ?, ?, ?, ?, ?, ? \rangle$  (*Fred enjoys sport on every day*)

- Intuitively a hypothesis  $h_i$  is more general than another  $h_j$  if
  - every instance that  $h_j$  classifies as positive  $h_i$  also classifies as positive, and
  - $h_i$  classifies instances as positive that  $h_j$  does not

E.g.

$h_1 = \langle \text{Sunny}, ?, ?, \text{Strong}, ?, ? \rangle$

$h_2 = \langle \text{Sunny}, ?, ?, ?, ?, ? \rangle$

$h_2$  is more general than  $h_1$

# Concept learning ...

- More formally, instance  $x$  **satisfies** hypothesis  $h$  iff  $h(x) = 1$ .


Define partial ordering relation *more\_general\_than\_or\_equal\_to* holding between two hypotheses  $h_i$  and  $h_j$  in terms of the sets of instances that satisfy them:

$h_i$  is *more\_general\_than\_or\_equal\_to*  $h_j$  iff every instance that satisfies  $h_j$  also satisfies  $h_i$   
or

**Definition:** Let  $h_i$  and  $h_j$  be boolean-valued functions defined over  $X$ .  $h_i$  is **more\_general\_than\_or\_equal\_to**  $h_j$  (written  $h_i \geq_g h_j$ ) if and only if

$$(\forall x \in X)[(h_j(x) = 1) \rightarrow (h_i(x) = 1)]$$

# Concept learning ...



- Note that in the example:

$$\begin{array}{ll}
 h_2 \geq_g h_1 & h_2 \geq_g h_3 \\
 h_1 \not\geq_g h_3 & h_3 \not\geq_g h_1
 \end{array}$$

## FIND-S:

### Finding a Maximally Specific Hypothesis

1. Initialize  $h$  to the most specific hypothesis in  $H$
2. For each positive training instance  $x$ 
  - For each attribute constraint  $a_i$  in  $h$ 
    - If the constraint  $a_i$  in  $h$  is satisfied by  $x$
    - Then do nothing
    - Else replace  $a_i$  in  $h$  by the next more general constraint that is satisfied by  $x$
3. Output hypothesis  $h$

## FIND-S ...

$x_1 = \langle \text{Sunny Warm Normal Strong Warm Same} \rangle, +$	$h_0 = \langle \emptyset, \emptyset, \emptyset, \emptyset, \emptyset, \emptyset \rangle$
$x_2 = \langle \text{Sunny Warm High Strong Warm Same} \rangle, +$	$h_1 = \langle \text{Sunny Warm Normal Strong Warm Same} \rangle$
$x_3 = \langle \text{Rainy Cold High Strong Warm Change} \rangle, -$	$h_2 = \langle \text{Sunny Warm ? Strong Warm Same} \rangle$
$x_4 = \langle \text{Sunny Warm High Strong Cool Change} \rangle, +$	$h_3 = \langle \text{Sunny Warm ? Strong Warm Same} \rangle$
	$h_4 = \langle \text{Sunny Warm ? Strong ? ?} \rangle$

Note: negative training instances completely ignored by FIND-S

## FIND-S ...

- For hypothesis spaces described by conjunctions of attribute constraints (e.g.  $H$  for *EnjoySport*), FIND-S is guaranteed to output the most specific hypothesis in  $H$  consistent with positive training examples
- FIND-S also guaranteed to output hypothesis consistent with negative examples provided
  - correct target concept is in  $H$
  - training examples are correct

# FIND-S ...

But, there are problems with FIND-S:

- may be multiple hypotheses consistent with the training data – FIND-S will find one, but give no indication of whether there may be others
- FIND-S always proposes maximally specific hypothesis – why prefer this to, e.g., maximally general?
- FIND-S has serious problems when training examples are inconsistent which frequently happens with noisy “real” data

## Version Spaces + Candidate Elimination

- One limitation of the **FIND-S** algorithm is that it outputs just one hypothesis consistent with the training data – there might be many.

To overcome this, introduce notion of **version space** and algorithms to compute it.

- A hypothesis  $h$  is **consistent** with a set of training examples  $D$  of target concept  $c$  if and only if  $h(x) = c(x)$  for each training example  $\langle x, c(x) \rangle$  in  $D$ .

$$\text{Consistent}(h, D) \equiv (\forall \langle x, c(x) \rangle \in D) h(x) = c(x)$$

- The **version space**,  $V_{S_H, D}$ , with respect to hypothesis space  $H$  and training examples  $D$ , is the subset of hypotheses from  $H$  consistent with all training examples in  $D$ .

$$V_{S_H, D} \equiv \{h \in H | \text{Consistent}(h, D)\}$$

# Version Spaces + Candidate Elimination

- Note difference between definitions of *consistent* and *satisfies*:
  - an example  $x$  *satisfies* hypothesis  $h$  when  $h(x) = 1$ , regardless of whether  $x$  is +ve or -ve example of target concept
  - an example  $x$  is *consistent* with hypothesis  $h$  iff  $h(x) = c(x)$

## The List-Then-Eliminate Algorithm

- Can represent version space by listing all members.
- Leads to **LIST-THEN-ELIMINATE** concept learning algorithm:

1.  $VersionSpace \leftarrow$  a list containing every hypothesis in  $H$
2. For each training example,  $\langle x, c(x) \rangle$   
remove from  $VersionSpace$  any hypothesis  $h$  for which  $h(x) \neq c(x)$
3. Output the list of hypotheses in  $VersionSpace$

- **LIST-THEN-ELIMINATE** works in principle, so long as version space is finite.
- However, since it requires exhaustive enumeration of all hypotheses in practice it is not feasible.



# The Candidate-Elimination Algorithm

- The **CANDIDATE-ELIMINATION** algorithm is similar to **LIST-THEN-ELIMINATE** algorithm but uses a more compact representation of version space.
  - represents version space by its most general and most specific members
- The **CANDIDATE-ELIMINATION** algorithm represents the version space by recording only the most general members ( $G$ ) and its most specific members ( $S$ )
  - other intermediate members in general-to-specific ordering can be generated as needed

# The Candidate-Elimination Algorithm

- The **General boundary**,  $G$ , of version space  $VS_{H,D}$  is the set of its maximally general members
- The **Specific boundary**,  $S$ , of version space  $VS_{H,D}$  is the set of its maximally specific members
- **Version Space Representation Theorem**  
Every member of the version space lies between these boundaries

$$VS_{H,D} = \{h \in H \mid (\exists s \in S)(\exists g \in G)(g \geq_g h \geq_g s)\}$$

where  $x \geq_g y$  means  $x$  is more general or equal to  $y$

# The Candidate-Elimination Algorithm

- Intuitively, **CANDIDATE-ELIMINATION** algorithm proceeds by
  - initialising  $G$  and  $S$  to the maximally general and maximally specific hypotheses in  $H$
  - considering each training example in turn and
    - \* using positive examples to drive the maximally specific boundary up
    - \* using negative examples to drive the maximally general boundary down

$G \leftarrow$  maximally general hypotheses in  $H$

$S \leftarrow$  maximally specific hypotheses in  $H$

For each training example  $d$ , do

- If  $d$  is a positive example
  - Remove from  $G$  any hypothesis inconsistent with  $d$
  - For each hypothesis  $s$  in  $S$  that is not consistent with  $d$ 
    - \* Remove  $s$  from  $S$
    - \* Add to  $S$  all minimal generalizations  $h$  of  $s$  such that
      1.  $h$  is consistent with  $d$ , and
      2. some member of  $G$  is more general than  $h$
    - \* Remove from  $S$  any hypothesis that is more general than another hypothesis in  $S$
- If  $d$  is a negative example
  - Remove from  $S$  any hypothesis inconsistent with  $d$
  - For each hypothesis  $g$  in  $G$  that is not consistent with  $d$ 
    - \* Remove  $g$  from  $G$
    - \* Add to  $G$  all minimal specializations  $h$  of  $g$  such that
      1.  $h$  is consistent with  $d$ , and
      2. some member of  $S$  is more specific than  $h$
    - \* Remove from  $G$  any hypothesis that is less general than another hypothesis in  $G$

# The Candidate-Elimination Algorithm

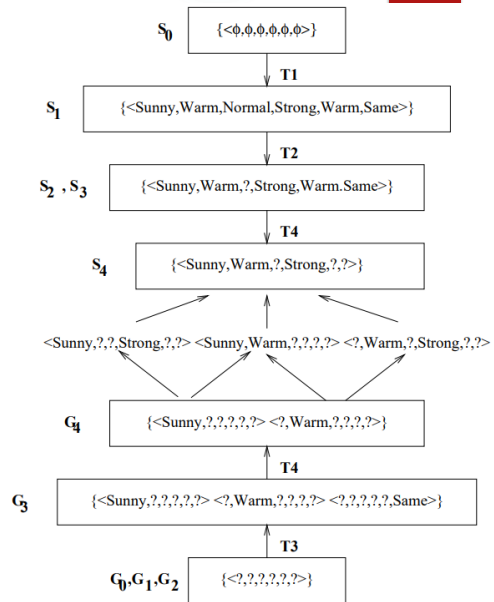
Training Examples:

T1: *<Sunny, Warm, Normal, Strong, Warm, Same>, Yes*

T2: *<Sunny, Warm, High, Strong, Warm, Same>, Yes*

T3: *<Rainy, Cold, High, Strong, Warm, Change>, No*

T4: *<Sunny, Warm, High, Strong, Cool, Change>, Yes*



# The Candidate-Elimination Algorithm

- Version space learned by **CANDIDATE-ELIMINATION** algorithm will converge towards correct hypothesis provided:

- no errors in training examples
- there is a hypothesis in  $H$  that describes target concept

In such cases algorithm may converge to empty version space

- If algorithm can request next training example (e.g. from teacher) can increase speed of convergence by requesting examples that split the version space
  - E.g. T5: *<Sunny, Warm, Normal, Light, Warm, Same>* satisfies 3 hypotheses in previous example
    - \* If T5 positive,  $S$  generalised, 3 hypotheses eliminated
    - \* If T5 negative,  $G$  specialised, 3 hypotheses eliminated

# The Candidate-Elimination Algorithm

- As noted, version space learned by **CANDIDATE-ELIMINATION** algorithm will converge towards correct hypothesis provided:
  - no errors in training examples
  - there is a hypothesis in  $H$  that describes target concept